

---

# **Implementation of a real-time semantic retrieval system.**

**David Lorenzo Alfaro**

**Jun 17, 2021**



## CONTENTS:

<b>1</b>	<b>Semantic search (<code>semantic_search.py</code>)</b>	<b>3</b>
<b>2</b>	<b>Lexical search (<code>lexical_search.py</code>)</b>	<b>7</b>
<b>3</b>	<b>Code utilities</b>	<b>11</b>
3.1	Utils ( <code>utils.py</code> ) . . . . .	11
3.2	Plotter ( <code>plotter.py</code> ) . . . . .	17
<b>4</b>	<b>Code used for experiments.</b>	<b>21</b>
4.1	Speedup-recall tradeoff depending on the number of trees used in AN- NOY( <code>experiment_annoy_ntrees.py</code> ) . . . . .	21
4.2	Evaluate text summarization using different values of $k$ top sentences ( <code>experiment_text_summarization.py</code> ) . . . . .	21
<b>5</b>	<b>Visualization of BERT embeddings.</b>	<b>23</b>
5.1	Introduction. . . . .	23
5.2	Visualization of kNN neighbours of BERT's pretrained embeddings. . . . .	23
5.3	Exploring BERT contextual representations. . . . .	26
<b>6</b>	<b>Exploratory data analysis and data preprocessing.</b>	<b>29</b>
6.1	Introduction. . . . .	29
6.2	Dataset exploration. . . . .	29
6.3	Remove useless features. . . . .	33
6.4	Integrate book overviews into the dataset. . . . .	33
6.5	Remove instances with invalid language codes. . . . .	35
6.6	Remove noisy data from book titles. . . . .	37
6.7	Remove noisy data from book overviews. . . . .	37
6.8	Annex. Code to perform data scraping. . . . .	38
<b>7</b>	<b>Indices and tables</b>	<b>41</b>
	<b>Python Module Index</b>	<b>43</b>
	<b>Index</b>	<b>45</b>



The increasingly overwhelming amount of available natural language motivates the pressing need to find efficient and reliable computational techniques capable of processing and analysing this type of data for the purpose of achieving human-like natural language understanding for a wide range of downstream tasks.

Over the last decade, Natural Language Processing (NLP) has seen impressively fast growth, primarily favoured by the increase in computational power and the progress on unsupervised learning in linguistics. Moreover, historically successful statistical language modeling techniques have been largely replaced by novel neural language modeling based on Deep Learning models, exhibiting an unprecedented level of natural language understanding, contributing to reduce the gap between human communication and computer understanding.

NLP is the key to solve many technological challenges. Among the large number of applications this field has, and since this dissertation has been entirely focused on the study of different strategies to encode salient information about natural language, the experimental part of this work is primarily devoted to the implementation of a semantic information retrieval system, delivering search latencies suitable for real-time similarity matching.

Of course, the examined unsupervised pretrained representations have been trained on humongous data sets, devoting to that end massive amounts of computational resources, something we do not have the access to. It is, however, not only a matter of computational power. Gathering, preparing and processing data for a model to be fine-tuned is no easy task and requires knowledge, to a great extent, of both the underlying theoretical motivations and the specific implementation. Consequently, our work is pretty much aligned with the mentality in which these strategies sit on: to directly use the pretrained models for downstream tasks.

This document contains the documentation for all experiments conducted throughout the dissertation, including a thoroughly description of the implementation of the information retrieval system and the notebooks devoted to the dataset analysis and preprocessing and the analysis and visualization of BERT contextual embeddings.



## SEMANTIC SEARCH (SEMANTIC\_SEARCH.PY)

This class implements a semantic textual information retrieval system.

Author: David Lorenzo Alfaro.

```
class semantic_search.SemanticSearch(corpus: pandas.core.frame.DataFrame,
                                     embeddings_cache_path: str = None,
                                     encoding_strategy='title_overview',
                                     pretrained_model='paraphrase-distilroberta-
                                     base-v2', model_max_seq_length=512, pre-
                                     trained_crossencoder: str = None, an-
                                     noy_index_cache_path: str = None, an-
                                     noy_n_trees=576, _annoy_embedding_size=768)
```

This class implements a semantic textual information retrieval system, which allows for advanced features like retrieve and re-rank and Approximate Nearest Neighbours search.

```
__init__(corpus: pandas.core.frame.DataFrame, embeddings_cache_path: str = None,
          encoding_strategy='title_overview', pretrained_model='paraphrase-distilroberta-
          base-v2', model_max_seq_length=512, pretrained_crossencoder: str = None, an-
          noy_index_cache_path: str = None, annoy_n_trees=576, _annoy_embedding_size=768)
```

Constructor for a *SemanticSearch* instance.

**Important:** if you are willing to load the embeddings or the ANNOY index from disk you do not need to tune their respective specific parameters (they will be overlooked).

### Parameters

- **corpus** (*DataFrame*) – Book corpus. It should, at least, have the following columns (with the very same names):
  - *gr\_book\_id*: book unique identifier.
  - *title*: book titles.
  - *authors*: book authors.
  - *overview*: book overview.
- **embeddings\_cache\_path** (*str, optional*) – Filepath to store the computed embeddings or load the embeddings from. Defaults to *None*.
- **encoding\_strategy** (*str, optional*) – Encoding strategy to use. The encoding strategy must be a string containing the names of the features to include into the input of the encoder, each of them separated by an underscore ('\_'). For example, if you were to use the title and the overview as the encoding strategy, *encoding\_strategy* must be either *title\_overview* or *overview\_title*. Defaults to 'title\_overview'.

- **pretrained\_model** (*str*, *optional*) – The model *id* of a predefined *Sentence-Transformer* hosted inside a model repo on *sbert.net*. Defaults to ‘paraphrase-distilroberta-base-v2’. Defaults to ‘paraphrase-distilroberta-base-v2’.
- **model\_max\_seq\_length** (*int*, *optional*) – Property to get the maximal input sequence length for the model. Longer inputs will be truncated. Defaults to 512.
- **pretrained\_crossencoder** (*str*, *optional*) – Any model name from Huggingface Models Repository that can be loaded with *AutoModel*. Defaults to None.
- **annoy\_index\_cache\_path** (*str*, *optional*) – Filepath to store an ANNOY index or load it from disk. Defaults to None.
- **annoy\_n\_trees** (*int*, *optional*) – Number of trees to use in the forest for ANNOY. Defaults to 576
- **\_annoy\_embedding\_size** (*int*, *optional*) – Size of the embeddings, required to compute the index. Defaults to 768

### **property** annoy\_index

Getter method for *annoy\_index*

**Returns** Current ANNOY index. If none, it attempts to create a new one with the optimal configuration (according to the experiments detailed in the dissertation document).

**Return type** AnnoyIndex

### **property** crossencoder

Getter method for *crossencoder*.

**Returns** Current pretrained Cross-Encoder. If none, it attempts to obtain the default one.

**Return type** CrossEncoder

**search** (*query*: *str*, *k*=5, *k\_biencoder*=20, *use\_annoy*=False, *reranking*=False)

Perform semantic search.

#### **Parameters**

- **query** (*str*) – Textual query.
- **k** (*int*, *optional*) – Number of most relevant documents to retrieve. When using exhaustive search, the value of *k* does not affect performance. Complexity using ANNOY and *k* ~ corpus length will be close to O(n). Defaults to 5
- **k\_biencoder** (*int*, *optional*) – If using retrieve and re-rank, number of documents to retrieve by the Bi-encoder and fed into the Cross-Encoder. The Cross-Encoder will return the *k* most relevant entries. *k\_biencoder* must be greater or equal to *k*. Defaults to 20.
- **use\_annoy** (*bool*, *optional*) – Use approximate search to reduce search time to approx O(log(n)). Defaults to False
- **reranking** (*bool*, *optional*) – Use retrieve and Re-Rank Pipeline, defaults to False

**search\_multiple** (*queries*: *list*, *write\_to*: *str* = None, *\*\*search\_options*)

Perform semantic search for several queries. Refer to the documentation of *search* method for further information.

#### **Parameters**

- **queries** (*list*) – Collection of textual queries.



- **write\_to** (*str*, *optional*) – Path of the file in which the results of the query will be written. If the file already exists, existing data will be overwritten. Defaults to None.

**setup\_annoy** (*index\_cache\_path*: *str* = None, *n\_trees*=576, *embedding\_size*=768)

Setup for ANNOY index. Use this method to:

- Use a precomputed ANNOY index located in *index\_cache\_path*.
- **Create a new ANNOY index and store it in *index\_cache\_path*.** if *index\_cache\_path* is None, the index will not be stored in disk.
- **Either way, the obtained ANNOY index will be used in future calls** to *search* and *search\_multiple* if approximate search is chosen.
- Previous ANNOY setup is replaced upon invoking this method.

**IMPORTANT:** if you are attempting to load an ANNOY index from disk, there is no need to tune the remaining parameters (i.e., *n\_trees* and *embedding\_size*)

#### Parameters

- **index\_cache\_path** (*str*, *optional*) – Filepath to store the obtained ANNOY index or filepath of a precomputed ANNOY index. By default is None: a new ANNOY index will be created with the indicated parameters.
- **n\_trees** (*int*, *optional*) – Number of trees to use in the forest for ANNOY, defaults to 576.
- **embedding\_size** (*int*, *optional*) – Size of the embeddings, required to compute the index. Defaults to 768

**setup\_crossencoder** (*pretrained\_crossencoder*='cross-encoder/stsb-distilroberta-base')

Set or update the Cross-Encoder to be used to re-rank the results retrieved by Bi-Encoder. We recommend using either 'cross-encoder/stsb-distilroberta-base' or any pretrained Cross-Encoder trained on MS MARCO dataset.

**Parameters pretrained\_crossencoder** (*str*, *optional*) – Any model name from Huggingface Models Repository that can be loaded with AutoModel. Defaults to 'cross-encoder/stsb-distilroberta-base'

**test\_annoy\_performance** (*queries*: *list*, *k*=5, *verbose*=True)

Utility to test the performance of ANNOY, considering the speedup with respect to exhaustive search and the recall.

#### Parameters

- **queries** (*list* or *array-like*) – Collection of textual queries.
- **k** (*int*, *optional*) – Top k elements to consider in the comparison.



## LEXICAL SEARCH (LEXICAL\_SEARCH.PY)

This class implements a textual literal information retrieval system.

Author: David Lorenzo Alfaro.

```
class lexical_search.TfIdfSearch(corpus: pandas.core.frame.DataFrame, vectors_cache_path:  
                                str = None, encoding_strategy='title_overview', pre-  
                                trained_biencoder: str = None, embeddings_cache_path:  
                                str = None, biencoder_max_seq_length=512, pre-  
                                trained_crossencoder: str = None)
```

This class implements a textual literal information retrieval system, based on TF-IDF which allows for advanced features like hybrid search, combining literal and dense search (retrieve and re-rank search pipeline).

```
__init__(corpus: pandas.core.frame.DataFrame, vectors_cache_path: str = None, encod-  
          ing_strategy='title_overview', pretrained_biencoder: str = None, embeddings_cache_path:  
          str = None, biencoder_max_seq_length=512, pretrained_crossencoder: str = None)
```

Constructor for a *TfIdfSearch* instance.

**Important:** if you are willing to load the embeddings or the vectors from disk you do not need to tune their respective specific parameters (they will be overlooked).

### Parameters

- **corpus** (*DataFrame*) – Book corpus. It should, at least, have the following columns (with the very same names):
  - *gr\_book\_id*: book unique identifier.
  - *title*: book titles.
  - *authors*: book authors.
  - *overview*: book overview.
- **vectors\_cache\_path** (*str, optional*) – Filepath to store the computed vectors or load the vectors from. Defaults to *None*.
- **encoding\_strategy** (*str, optional*) – Encoding strategy to use. The encoding strategy must be a string containing the names of the features to include into the input of the encoder, each of them separated by an underscore ('\_'). For example, if you were to use the title and the overview as the encoding strategy, *encoding\_strategy* must be either *title\_overview* or *overview\_title*. Defaults to 'title\_overview'.
- **pretrained\_biencoder** (*str, optional*) – The model *id* of a predefined *SentenceTransformer* hosted inside a model repo on *sbnet.net*. Defaults to 'paraphrase-distilroberta-base-v2'. Defaults to *None*.
- **embeddings\_cache\_path** (*str, optional*) – Filepath to store the computed embeddings or load the embeddings from. Defaults to *None*.

- **biencoder\_max\_seq\_length** (*int, optional*) – Property to get the maximal input sequence length for the model. Longer inputs will be truncated. Defaults to 512.
- **pretrained\_crossencoder** (*str, optional*) – Any model name from Huggingface Models Repository that can be loaded with AutoModel. Defaults to None.

### **property biencoder**

Getter method for *biencoder*.

**Returns** Current pretrained Bi-Encoder. If none, it attempts to obtain the default one.

**Return type** SentenceTransformer

### **property crossencoder**

Getter method for *crossencoder*.

**Returns** Current pretrained Cross-Encoder. If none, it attempts to obtain the default one.

**Return type** CrossEncoder

**search** (*query: str, k=5, k\_lexical=20, reranking\_strategy: Literal[crossencoder, biencoder] = None*)  
Perform TF-IDF lexical search.

#### **Parameters**

- **query** (*str*) – Textual query.
- **k** (*int, optional*) – Number of most relevant documents to retrieve.
- **k\_lexical** (*int, optional*) – If using retrieve and re-rank, number of documents to retrieve by lexical search and re-ranked by any re-ranking strategy. Re-Ranker will return the *k* most relevant entries. *k\_lexical* must be greater or equal to *k*. Defaults to 20.
- **reranking\_strategy** (*Literal['crossencoder', 'biencoder'], optional*) – Re-ranking strategy to use. Defaults to None

**Raises ValueError** – Raise *ValueError* if *reranking\_strategy* takes an illegal value.

**search\_multiple** (*queries: list, write\_to: str = None, \*\*search\_options*)

Perform lexical search for several queries. Refer to the documentation of *search* method for further information.

#### **Parameters**

- **queries** (*list*) – Collection of textual queries.
- **write\_to** (*str, optional*) – Path of the file in which the results of the query will be written. If the file already exists, existing data will be overwritten. Defaults to None.

**setup\_biencoder** (*pretrained\_biencoder='paraphrase-distilroberta-base-v2', max\_seq\_length=512*)

Set or update the Bi-Encoder to be used to re-rank the results retrieved by TF-IDF search.

#### **Parameters**

- **pretrained\_model** (*str, optional*) – The model *id* of a predefined *Sentence-Transformer* hosted inside a model repo on sbert.net. Defaults to 'paraphrase-distilroberta-base-v2'. Defaults to 'paraphrase-distilroberta-base-v2'.
- **max\_seq\_length** (*int, optional*) – Property to get the maximal input sequence length for the model. Longer inputs will be truncated. Defaults to 512.

**setup\_crossencoder** (*pretrained\_crossencoder='cross-encoder/stsb-distilroberta-base'*)

Set or update the Cross-Encoder to be used to re-rank the results retrieved by TF-IDF search. We recommend using either 'cross-encoder/stsb-distilroberta-base' or any pretrained Cross-Encoder trained on MS MARCO dataset.

**Parameters** `pretrained_crossencoder` (*str*, *optional*) – Any model name from Huggingface Models Repository that can be loaded with `AutoModel`. Defaults to ‘cross-encoder/stsb-distilroberta-base’



## CODE UTILITIES

Documentation for the *code\_utils* Python package.

### 3.1 Utils (*utils.py*)

Set of miscellaneous utilities used across the implementation.

Author: David Lorenzo Alfaro.

```
code_utils.utils.append_overviews_to_data(df_overviews: pandas.core.frame.DataFrame,
                                          df_data: pandas.core.frame.DataFrame,
                                          overview_index: str = 'gr_book_id', data_index:
                                          str = 'gr_book_id', merge_option='right')
```

Merge *df\_overviews* with *df\_data* using column identifiers *overview\_index* and *data\_index*, respectively. We allow books with no overviews, hence *right join* is the most suitable operation.

#### Parameters

- **df\_overviews** (*pd.DataFrame*) – Book overviews.
- **df\_data** (*pd.DataFrame*) – Book data (e.g., title, authors, etc.)
- **overview\_index** (*str*, *optional*) – Column or index level names to join on in the left DataFrame, defaults to 'gr\_book\_id'.
- **data\_index** (*str*, *optional*) – Column or index level names to join on in the right DataFrame, defaults to 'gr\_book\_id'.
- **merge\_option** (*str*, *optional*) – Type of merge operation, to be performed can be one of {'left', 'right', 'outer', 'inner'}, to 'right'.

**Returns** A DataFrame of the two merged objects.

**Return type** *pd.DataFrame*

```
code_utils.utils.clean_book_title(title: str, remove_quotation_marks=True, re-
                                move_saga_info=False, remove_saga_number=True)
```

Applies several transformations to a book title to remove noisy data that can potentially affect the performance of the embedding strategies.

#### Parameters

- **title** (*str*) – Book title in plain text.
- **remove\_quotation\_marks** (*bool*, *optional*) – If set to True, attempts to remove the quotation marks enclosing the book title, to True.

- **remove\_saga\_info** (*bool, optional*) – If set to True, attempts to remove information concerning the book saga, defaults to False.
- **remove\_saga\_number** (*bool, optional*) – If set to True, attempts to remove the saga number, defaults to True.

**Returns** Processed book title.

**Return type** str

`code_utils.utils.clean_overview` (*overview: str*)

Applies several transformations to a book overview to remove noisy data that can potentially affect the performance of the embedding strategies. One must be careful when applying transformations to the whole corpus because the odds for negative side-effects are high. Here, we attempt to solve some of the problems spotted that, in our tests, should not have any noticeable negative effect on any book overview.

**Parameters** **overview** (*str*) – Book overview in plain text.

**Returns** Processed book overview.

**Return type** str

`code_utils.utils.compute_avg_wordpiece_tokens` (*corpus: list, tokenizer=None*)

Compute the average number of WordPiece tokens in a list of documents, *corpus*.

**Parameters**

- **corpus** (*list or array-like*) – List of textual documents.
- **tokenizer** (*BertTokenizer, optional*) – Instance of *BertTokenizer* class. If *None*, it loads the predefined tokenizer of ‘bert-base-uncased’.

**Returns** Average number of WordPiece tokens of the documents in *corpus*.

**Return type** int

`code_utils.utils.fix_punctuation` (*overview: str*)

Attempts to fix some of the identified punctuation issues present in the book overviews.

**It is a common issue to find overviews with the following punctuation flaw:**

- “[...] word.Word [...]”
- “[...] word!Word [...]”
- “[...] word?Word [...]”

That is to say, spacing after periods, exclamation and question marks is not correctly applied. This lead to some issues when splitting the text into sentences, specially using the NLTK library. Furthermore, it may have other adverse effects on the embedding process (e.g., due to faulty tokenization).

**Parameters** **overview** (*str*) – Book overview in plain text.

**Returns** str

**Return type** Book overview without the identified punctuation flaws.

`code_utils.utils.generate_dataframe_from_sparse_txts` (*base\_dir,*  
*path\_standard\_format=False,*  
*out\_filename=None*)

Generates a dataframe from all *txt* files located in *base\_dir*. The dataframe features two columns: *gr\_book\_id*, an identifier that is retrieved from the name of each *txt* file, and *overview*, containing all information included in the *txt* file.

**Parameters**

- **base\_dir** (*str*) – Directory in which the *txt* files for the book overviews are located.



- **path\_standard\_format** (*bool, optional*) – Indicates whether the path follows the standard format (backslash separator) or the slash separator, defaults to False.
- **out\_filename** (*str, optional*) – Path for the output file, defaults to None.

**Returns** Returns a dataframe from all txt files located in *base\_dir*.

**Return type** `pd.DataFrame`

`code_utils.utils.get_bert_model (transformer='bert-base-uncased')`

Get an instance of the class *BertModel* for the transformer *trasformer*.

Wrapper function of the HuggingFace Transformer's *BertModel* function: *from\_pretrained*.

**Parameters** **transformer** (*str, optional*) – The model *id* of a predefined tokenizer hosted inside a model repo on huggingface.co. Defaults to 'bert-base-uncased'

:return Instance of *BertTokenizer* class. :rtype: *BertTokenizer*

`code_utils.utils.get_bert_tokenizer (transformer='bert-base-uncased')`

Get an instance of the class *BertTokenizer* for the transformer *trasformer*.

Wrapper function of the HuggingFace Transformer's *BertTokenizer* function: *from\_pretrained*.

**Parameters** **transformer** (*str, optional*) – The model *id* of a predefined tokenizer hosted inside a model repo on huggingface.co. Defaults to 'bert-base-uncased'

:return Instance of *BertTokenizer* class. :rtype: *BertTokenizer*

`code_utils.utils.get_crossencoder (crossencoder='cross-encoder/stsb-distilroberta-base',  
**kwargs)`

Wrapper function to get a *CrossEncoder*.

**Parameters** **crossencoder** (*str, optional*) – Any model name from Huggingface Models repository that can be loaded with *AutoModel*. Defaults to 'cross-encoder/stsb-distilroberta-base'.

**Returns** a *CrossEncoder* that takes exactly two sentences/texts as input and predicts a score for this sentence pair. It can for example predict the similarity of the sentence pair on a scale of 0 ... 1.

**Return type** *CrossEncoder*

`code_utils.utils.get_dir_files_content (directory: str, file_extension='.txt')`

Get the list of files in *directory* with extension *file\_extension*.

**Parameters**

- **directory** (*str*) – Directory in which the files are located. Recursive search is not allowed.
- **file\_extension** (*str, optional*) – File extension, defaults to '.txt'

**Returns** List of tuples (filename, file content).

**Return type** *list*

`code_utils.utils.get_file_id_and_content (filepath)`

Returns all textual content in 'filepath'

**Parameters** **filepath** – Path of the text file to read.

**Returns** Content of the file.

**Return type** *str*

`code_utils.utils.get_sentence_transformer` (*transformer='paraphrase-distilroberta-base-v2', max\_seq\_length=None, \*\*kwargs*)

Wrapper function to get a *SentenceTransformer*.

### Parameters

- **transformer** (*str, optional*) – The model *id* of a predefined *SentenceTransformer* hosted inside a model repo on *sbert.net*. Defaults to ‘paraphrase-distilroberta-base-v2’.
- **max\_seq\_length** (*int*) – Property to get the maximal input sequence length for the model. Longer inputs will be truncated. Defaults to None.

**Returns** a *SentenceTransformer* model that can be used to map sentences / text to embeddings.

**Return type** *SentenceTransformer*

`code_utils.utils.get_tokenized_text` (*text: str, tokenizer=None*)

Get the list of WordPiece tokens in *text*

### Parameters

- **text** (*str*) – Text to tokenize
- **tokenizer** (*BertTokenizer, optional*) – Instance of *BertTokenizer* class. If *None*, it loads the predefined tokenizer of ‘bert-base-uncased’.

**Returns** list of WordPiece tokens.

**Return type** list

`code_utils.utils.get_top_k_sentences` (*document: str, k=5, embedder=None, spacy\_model=None, preserve\_order=True*)

Get the top *k* most meaningful sentences of *document*.

### Parameters

- **document** (*str*) – a document in plain text
- **k** (*int, optional*) – Top *k* sentences to return, defaults to 5
- **embedder** (*SentenceTransformer, optional*) – Instance of *SentenceTransformer*. If *None*, it loads the default pretrained sentence transformer model. Defaults to *None*.
- **spacy\_model** (*str, optional*) – Name of a spacy pretrained tokenizer model. If *None*, it loads the default model. Defaults to *None*.
- **preserve\_order** (*bool, optional*) – Preserve the order of the top *k* sentences with respect to the original document to conserve spatial dependencies between sentences. Defaults to *True*.

**Returns** Top *k* most meaningful sentences of *document*.

**Return type** list

`code_utils.utils.load_corpus` (*path\_corpus: str, sep=',', \*\*kwargs*)

Wrapper method of Pandas *read\_csv* function to load book corpus.

### Parameters

- **path\_corpus** (*str*) – Filepath to the corpus.
- **sep** (*str, optional*) – Delimiter to use, defaults to ‘,’

**Returns** Corpus DataFrame

**Return type** DataFrame

```
code_utils.utils.load_embeddings(filename: str, return_dict_values=True)
```

Utility to load embeddings (and other optional stored values) from disk using *pickle*.

#### Parameters

- **filename** (*str*) – Filename of the file to be loaded.
- **return\_dict\_values** (*bool, optional*) – If set to True, returns the values just the values of the dictionary containing all stored data, defaults to True.

**Returns** Loaded data

```
code_utils.utils.makedir(path: str, remove_filename=False, recursive=True, exist_ok=True)
```

Creates directory from path if not exists.

#### Parameters

- **path** (*str*) – Path of the directory to be created.
- **remove\_filename** (*bool, optional*) – If set to True, it attempts to remove the filename from the path, defaults to False
- **recursive** (*bool, optional*) – Creates directories recursively (i.e., create necessary subdirectories if necessary), defaults to True
- **exist\_ok** (*bool, optional*) – is set to False, arises an error if *path* directory exists, defaults to True

```
code_utils.utils.prepare_input_encoder(encoding_strategy: str, corpus: pandas.core.frame.DataFrame, return_input_encoder=True)
```

Formats the input to the encoder using the features indicated in *encoding\_strategy*. If *encoding\_strategy* takes a wrong value this method is likely to fail. Current supported features are ‘title’, ‘authors’ and ‘overview’.  
: param str encoding\_strategy: The encoding strategy must be a string containing the names of the features to include into the input of the encoder, each of them separated by an underscore (‘\_’). For example, if you were to use the title and the overview as the encoding strategy, *encoding\_strategy* must be either *title\_overview* or *overview\_title*.  
: param str path\_df: Path in which the dataframe is located.  
: param return\_input\_encoder: Return just the collection of inputs to the encoder, defaults to True  
: type return\_input\_encoder: bool, optional  
: return: If *return\_input\_encoder*, returns the collection of inputs to the encoder. Otherwise, it returns Dataframe including a new column *input\_encoder* with the format indicated in *encoding\_strategy*  
: rtype: Dataframe

```
code_utils.utils.remove_filename_from_path(out_filename: str, path_standard_format=False)
```

Attempts to remove filename from the provided path.

#### Parameters

- **out\_filename** (*str*) – Filepath.
- **path\_standard\_format** (*bool, optional*) – Indicates whether the path follows the standard format (backslash separator) or the slash separator, defaults to False.

**Returns** The directory excluding the filename.

**Return type** str

```
code_utils.utils.split_text_into_sentences_nltk(text: str)
```

Splits text into sentences using the NLTK library.

**Parameters** **text** (*str*) – Text to be splitted into sentences.

**Returns** List of sentences in *text*.

**Return type** list

`code_utils.utils.split_text_into_sentences_spacy(text, spacy_model='en_core_web_sm')`

Splits text into sentences using the Spacy library. SpaCy builds a syntactic tree for each sentence, a robust method that yields more statistical information about the text than NLTK. It performs substantially better than NLTK when using not polished text.

### Parameters

- **text** (*str*) – Text to be splitted into sentences.
- **spacy\_model** (*str*, *optional*) – Name of the spacy pretrained model used to split text into sentences, defaults to 'en\_core\_web\_sm'.

**Returns** List of sentences in *text*.

**Return type** list

`code_utils.utils.store_embeddings(corpus_embeddings, out_filename='embeddings.pkl', protocol=5, **kwargs)`

Utility to dump embeddings (and other optional values indicated in the keyword arguments) to disk using *pickle*.

### Parameters

- **corpus\_embeddings** – Tensor type data structure containing the embeddings for the corpus.
- **out\_filename** (*str*, *optional*) – Path for the output file, defaults to 'embeddings.pkl'.
- **protocol** – Protocol used for *pickle*, defaults to *pickle.HIGHEST\_PROTOCOL*.

`code_utils.utils.summarize_corpus_overviews(corpus_overviews: list, top_k=5, embedder=None, spacy_model=None, **kwargs)`

Apply unsupervised Text Summarization techniques to obtain representations for the most meaningful sentences for each document in *corpus\_overviews*.

### Parameters

- **corpus\_overviews** (*list* or *array-like*) – Book overviews.
- **top\_k** (*int*, *optional*) – Number of sentences that will have each overview. Defaults to 5.
- **embedder** (*SentenceTransformer*, *optional*) – Instance of *SentenceTransformer*. If *None*, it loads the default pretrained sentence transformer model. Defaults to *None*.
- **spacy\_model** (*str*, *optional*) – Name of a spacy pretrained tokenizer model. If *None*, it loads the default model. Defaults to *None*.

**Returns** Summarized overviews with at most *top\_k* sentences.

**Return type** list

`code_utils.utils.topk_cos_sim(query_embedding: torch.Tensor, embeddings: torch.Tensor, top_k: int)`

Get the indices and the cosine similarity score of the *top\_k* most similar embeddings to *query\_embedding* in *embeddings*.

### Parameters

- **query\_embedding** (*torch.Tensor*) – Query embedding.
- **embeddings** (*torch.Tensor*) – Corpus embeddings.
- **top\_k** (*int*) – Top k most similar to retrieve according to cosine similarity score.

**Returns** List of scores and list of indexes of the top k results.

**Return type** (list, list)

## 3.2 Plotter (plotter.py)

Set of plotting utilities used across the implementation.

Author: David Lorenzo Alfaro.

```
code_utils.plotter.apply_dimensionality_reduction(word_embeddings, n_components:
                                                    int, reduction_strategy='pca',
                                                    random_state=None,
                                                    _tsne_perplexity=5,
                                                    _tsne_learning_rate=10,
                                                    _tsne_n_iter=3000)
```

Apply dimensionality reduction on words embeddings using reduction techniques like the Principal Component Analysis (PCA) and the T-distributed Stochastic Neighbour Embedding (t-SNE). The default values for the perplexity, learning rate and number of iterations have been empirically tuned to those that produced acceptable results consistently.

### Parameters

- **word\_embeddings** (*list or array-like*) – Collection of word embeddings.
- **n\_components** (*int*) – Dimension of the new embedded space (e.g., 2 for 2D visualization, 3 for 3D visualization).
- **random\_state** (*int, optional*) – Random state for dimensionality reduction techniques, defaults to None
- **reduction\_strategy** (*str, optional*) – Reduction strategy to choose. Can be either 'pca' or 'tsne'. Defaults to 'pca'
- **\_tsne\_perplexity** (*int, optional*) – The perplexity is related to the number of nearest neighbors that is used in other manifold learning algorithms. Defaults to 5.
- **\_tsne\_learning\_rate** (*int, optional*) – The learning rate for t-SNE is usually in the range [10.0, 1000.0]. Defaults to 10.
- **\_tsne\_n\_iter** (*int, optional*) – Maximum number of iterations without progress to abort optimization process, defaults to 3000.

**Raises** **ValueError** – Raise *ValueError* if *reduction\_stragy* takes an ilegal value.

**Returns** Data in the new embedded space.

**Return type** NumPy array

```
code_utils.plotter.display_embeddings_scatterplot_2D(word_embeddings,
                                                    color_text=None,          ran-
                                                    dom_state=None,          re-
                                                    duction_strategy='pca',
                                                    _graph_showlegend=True,
                                                    tsne_params={},          **scat-
                                                    ter_params)
```

Visualize BERT embeddings in a 2D scatterplot using dimensionality reduction techniques like Principal Component Analysis (PCA) and the T-distributed Stochastic Neighbour Embedding (t-SNE).

### Parameters

- **word\_embeddings** (*list or array-like*) – Collection of word embeddings.
- **color\_text** (*str, optional*) – Label for plotly scatterplot *color* attribute. Defaults to None
- **random\_state** (*int, optional*) – Random state for dimensionality reduction techniques, defaults to None
- **reduction\_strategy** (*str, optional*) – Reduction strategy to choose. Can be either 'pca' or 'tsne'. Defaults to 'pca'
- **\_graph\_showlegend** (*bool, optional*) – Show legend, defaults to True
- **tsne\_params** – Additional parameters for t-SNE, defaults to {}

```
code_utils.plotter.display_embeddings_scatterplot_3D(word_embeddings,
                                                    color_text: str = None,
                                                    random_state=None,      re-
                                                    duction_strategy='pca',
                                                    _graph_showlegend=True,
                                                    tsne_params={},          **scat-
                                                    ter_params)
```

Visualize BERT embeddings in a 3D scatterplot using dimensionality reduction techniques like Principal Component Analysis (PCA) and the T-distributed Stochastic Neighbour Embedding (t-SNE).

### Parameters

- **word\_embeddings** (*list or array-like*) – Collection of word embeddings.
- **color\_text** (*str, optional*) – Label for plotly scatterplot *color* attribute. Defaults to None
- **random\_state** (*int, optional*) – Random state for dimensionality reduction techniques, defaults to None
- **reduction\_strategy** (*str, optional*) – Reduction strategy to choose. Can be either 'pca' or 'tsne'. Defaults to 'pca'
- **\_graph\_showlegend** (*bool, optional*) – Show legend, defaults to True
- **tsne\_params** – Additional parameters for t-SNE, defaults to {}

```
code_utils.plotter.histogram_embeddings_nn(data: list)
```

Plot histogram for the nearest neighbors of an embedding.

**Parameters** *data* (*list or array-like*) – Two dimensional array containing a collection of similar words (first component), list of similarity scores (second component), and a list of labels (third component).

```
code_utils.plotter.plot_bert_embeddings_nn(model, vocab: dict, input_words, k=5, dis-
                                          play_option='3d', reduction_strategy='pca',
                                          random_state=None)
```

Visualize the *k* most similar words in *vocab* in the 2D or 3D embedding space. (Disclaimer: sorry about poor code readability).

### Parameters

- **model** (*BertModel*) – Bert pretrained model.
- **vocab** (*dict*) – Tokenizer vocabulary.
- **input\_words** (*Iterable*) – Words, the KNN of which are to be calculated and displayed.
- **k** (*int, optional*) – number of similar words to visualize. By default, 5

- **display\_option** (*str*, *optional*) – Visualize BERT embeddings either in ‘2d’ or ‘3d’, defaults to ‘3d’
- **reduction\_strategy** (*str*, *optional*) – Reduction strategy to choose. Can be either ‘pca’ or ‘tsne’. Defaults to ‘pca’
- **random\_state** (*int*, *optional*) – Random state for dimensionality reduction techniques, defaults to None

**Raises** **ValueError** – Raise *ValueError* if *display\_option* takes an ilegal value.

`code_utils.plotter.plot_heatmap` (*data*, *color\_continuous\_scale*: *str* = None, *\*\*kwargs*)

Plot heatmap. Wrapper of the *imshow* function.

#### Parameters

- **data** (*array-like*) – 2D data to be plotted.
- **color\_continuous\_scale** (*str*, *optional*) – Colour scale to be used in the heatmap, defaults to None

`code_utils.plotter.plot_heatmap_embeddings` (*model*, *tokenizer*, *data*: *pandas.core.frame.DataFrame*, *polysemous\_word*: *str*, *color\_continuous\_scale*: *str* = None)

Plot heatmap of the cosine similarity of all different contextual embeddings of *polysemous\_word* in *data*. This experiment is explained in the “Exploring BERT contextual representations” section of the dissertation. (Disclaimer: sorry about poor code readability).

#### Parameters

- **model** (*BertModel*) – Bert pretrained model.
- **tokenizer** (*BertTokenizer*) – Bert precomputed tokenizer.
- **data** (*pd.DataFrame*) – Test data for WSD evaluation.
- **polysemous\_word** (*str*) – Polysemous word.
- **color\_continuous\_scale** (*str*, *optional*) – Colour scale to be used in the heatmap, defaults to None

**Returns** The collection of the different contextual embeddings, along with the labels.

`code_utils.plotter.plot_scatter_with_secondary_y_axis` (*x*, *y*, *y2*, *fig\_title*=“, *x\_title*=“, *y\_title*=“, *y2\_title*=“)

Plot scatter plot with secondary y axis.

`code_utils.plotter.write_embeddings_to_disk` (*model*, *vocab*: *dict*, *out\_dir*=‘runs/bert\_embeddings’, *write\_word\_embeddings*=True, *write\_position\_embeddings*=False, *write\_type\_embeddings*=False)

**Utility to write embeddings weights to disk that can be loaded with** TensorBoard to visualize the embeddings.

#### Parameters

- **model** (*BertModel*) – Bert pretrained model.
- **vocab** (*dict*) – Tokenizer vocabulary.
- **out\_dir** (*str*, *optional*) – Directory to write the embeddings, defaults to ‘runs/bert\_embeddings’.

- **write\_word\_embeddings** (*bool*, *optional*) – Write word embeddings to disk, defaults to True.
- **write\_position\_embeddings** (*bool*, *optional*) – Write position embeddings to disk, defaults to False.
- **write\_type\_embeddings** (*bool*, *optional*) – Write token type embeddings to disk, defaults to False.



## CODE USED FOR EXPERIMENTS.

Documentation for the experiments described in the dissertation.

### 4.1 Speedup-recall tradeoff depending on the number of trees used in ANNOY(`experiment_annoy_ntrees.py`)

Experiment to test speedup-recall tradeoff depending on the number of trees used in ANNOY.

Author: David Lorenzo Alfaro.

`experiment_annoy_ntrees.evaluate_n_trees` (*queries: list, search\_alg, k=5, verbose=True*)

Utility used to evaluate the speedup-recall tradeoff of ANNOY as the number of trees increases.

#### Parameters

- **queries** (*list or array-like*) – Collection of textual queries.
- **search\_alg** (`SemanticSearch`) – Semantic search object.
- **k** (*int, optional*) – Top k elements to consider in the comparison, defaults to 5
- **verbose** (*bool, optional*) – Verbosity mode, defaults to True

### 4.2 Evaluate text summarization using different values of *k* top sentences (`experiment_text_summarization.py`)

Experiment to evaluate text summarization using different values of *k* top sentences.

Author: David Lorenzo Alfaro.

`experiment_text_summarization.evaluate_summarization_candidates` (*corpus\_overviews: list, candidates, embedder=None, spacy\_model=None, tokenizer=None, \*\*kwargs*)

Get an array of reduction rates and number of word pieces for a set of candidate number of sentences used to summarize each book overview (must be a list or array-like of integers).

#### Parameters

- **corpus\_overviews** (*list or array-like*) – Book overviews.
- **candidates** (*Iterable*) – List of number of candidates to evaluate.
- **embedder** (*SentenceTransformer, optional*) – Instance of *SentenceTransformer*. If *None*, it loads the default pretrained sentence transformer model. Defaults to *None*.
- **spacy\_model** (*str, optional*) – Name of a spacy pretrained tokenizer model. If *None*, it loads the default model. Defaults to *None*.
- **tokenizer** (*BertTokenizer, optional*) – Instance of *BertTokenizer* class. If *None*, it loads the predefined tokenizer of ‘bert-base-uncased’.

**Returns** Summarized overviews with at most *candidates* sentences.

**Return type** list

## VISUALIZATION OF BERT EMBEDDINGS.

As part of the undergraduate dissertation: Similarity Measures in Natural Language Processing based on Deep Learning Models.

David Lorenzo Alfaro

### 5.1 Introduction.

BERT embeddings have a dimensionality of 768. As studied, linguistic features are encoded along the representation of the word in a distributed fashion. Whilst embeddings work particularly well for computer systems to handle semantic and syntactic information about natural language, if humans were to visualize sequences of 768 floating point values, observing any valuable information would be an extremely difficult task. Furthermore, when handling natural language, the most common scenario involves dealing with sequences of words, thus posing neural NLP as a high-dimensionality problem. Consequently, the great majority of visualization techniques for embeddings involves dimensionality reduction.

The most common dimensionality reduction techniques are the Principal Component Analysis (PCA) and the T-distributed Stochastic Neighbour Embedding (t-SNE). The former is usually preferred because it can be used as a black box. That is, since it is a non-parametric method, it is not usually required to know the mathematics behind PCA.

The following sections provide a general overview of the different experiments carried out to visualize BERT embeddings. To that end, the increasingly popular open source graphing library Plotly (<https://plotly.com/>) is used. Plotly is a simple but expressive capable library that enables creating interactive charts and maps, a pretty much desired characteristic for the visualization tasks hereinafter described.

### 5.2 Visualization of kNN neighbours of BERT's pretrained embeddings.

The first test consist in gathering the WordPiece token embeddings weights learned through training of a BERT model for some query words to then find the k nearest neighbours to each word in the query. Let us first load all necessary libraries and modules.

```
[1]: import numpy as np
import pandas as pd
from code_utils.paths import *
# Plotter module contain a bunch of utilities for
# BERT embeddings visualization.
from code_utils import utils, plotter
```

For the sake of simplicity, we decided to use BERT's base uncased pretrained model. This model has way less parameters and weights, hence is lighter. We have not run this experiment on any pretrained model based on BERT Large architecture, albeit differences in outcomes, if any, must be negligible.

```
[2]: pretrained_model = 'bert-base-uncased'

# Get BERT pretrained model.
model = utils.get_bert_model(pretrained_model)

# Get BERT precomputed tokenizer.
tokenizer = utils.get_bert_tokenizer(pretrained_model)

# Get tokenizer vocabulary.
vocab = tokenizer.vocab

# To guarantee experiment reproducibility
random_state = 0
```

It is worth mentioning that the way in which experiment has been run is not appropriate for BERT for several reasons. First and foremost, BERT derives context-sensitive subword representations. For BERT to actually derive the proper embedding for a specific token, some context (i.e., a sentence containing that word) must be provided. In the approach we are following, comparisons among representations make no use of context, which would just be fine for context-independent embeddings such as those derived by techniques like GloVe or Word2vec. Second, since no actual embeddings are being computed (i.e., the learned weights for each token in BERT's vocab are used), all query words do necessarily need to belong to the BERT's model vocab, hence noticeably constraining the range of possible words to use in the query. Under these assumptions, we consider this experiment, albeit not entirely correct, can help the reader to have a better understanding of some of the concepts previously discussed, especially those concerning the linguistic regularities in continuous space word representations.

The process is straightforward. Given a subset of BERT's vocabulary tokens, get the learned embeddings and compute the  $k$  nearest neighbours to each query word. The set of neighbour candidates is the complete vocabulary used in BERT (30,000 tokens). Since it is a sine qua non condition that the query tokens belong to BERT's vocab, the nearest neighbour for each query token is necessarily the token itself. We found that the easiest and computationally cheapest way to prevent this from happening is to get the  $k + 1$  nearest neighbours to then discard those that are incorrect.

Afterwards, the proper embeddings for the kNN of all query tokens are fed into one dimensionality reduction technique. Since it is our will that the output of either PCA or t-SNE are visually interpretable, the embeddings, originally in the 768-dimensional space, are compressed into either 2-dimensional or 3-dimensional representations, attempting to preserve some of the meaningful information.

For further details, check the documentation for `plot_bert_embeddings_nn` method.

In the dissertation we discussed that Word2vec embeddings were able to capture features like the notion of gender or the syntactic singular/plural relation. For this first query, we will be using again the word queen.

```
[3]: query = 'queen'

# Format query words.
input_words = query.replace(' ', '').lower().split(',')

k=8

# Invoke method to plot PCA 2D projection of the embeddings
# and a histogram for all nearest neighbours of the query
# word embeddings.
plotter.plot_bert_embeddings_nn(model,
```

(continues on next page)

(continued from previous page)

```

vocab,
input_words,
k=k,
reduction_strategy='pca',
display_option='2d',
random_state=random_state)

```

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html

As it can be seen, the eight closest words to queen are, in order, king, queens, princess, empress, prince, duchess, countess and monarch. The results are astoundingly positive: all neighbours are reasonably related, either syntactically (e.g., queens) or semantically (e.g., king). Furthermore, the similarity scores obtained denote strong relatedness between the similar and query words.

It is also worth noting that, excepting king and prince (which, on their own are fairly related to queen), all remaining nearest neighbours are royal titles held by women, which hints BERT embeddings capability to capture the notion of gender. Moreover, the words empress, countess, and duchess seem to be closer in space, probably due to syntactic similarities among them (i.e., they all are suffixed with “ess”).

Let us now repeat the experiment using a set of queries

```

[4]: queries = 'stupid, queen, wizard, spain, brother'

# Format query words.
input_words = queries.replace(' ', '').lower().split(',')

k=5

# Invoke method to plot T-SN 2D projection of the embeddings
# and a histogram for all nearest neighbours of the query
# word embeddings.
plotter.plot_bert_embeddings_nn(model,
                                vocab,
                                input_words,
                                k=k,
                                reduction_strategy='t-sne',
                                display_option='2d',
                                random_state=random_state)

```

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html

As hinted by the results, each query word, along with its most similar words in terms of cosine-similarity, conform a cluster. Unless query words are highly correlated, this is to expect, being as cosine-similarity is a measurement of distance in the d-dimensional space. Consistently with the highlighted observations for the previous experiment,

all nearest neighbours to the query words are intimately connected syntactically and/or semantically, with the cosine similarity scores indicating strong relatedness.

Generally, some of the words populating the nearest neighbour set of each query word are synonyms (e.g., wizard and magician, stupid and dumb). Also, for the query word Spain, some of the most similar words are countries as well (e.g., Portugal, Italy). BERT embeddings showcase, once more, its capability to learn fine-grained features about natural language. For the query word brother some of the similar words are terms that belong to the lineal kinship system used in the English-speaking world (e.g., son, father).

It is also worth noting that, according to empirical observations, t-SNE dimensionality reduction technique is encouraged over PCA as the number of query words increase, being as it has reportedly produced better defined clusters, which favours the proper visualization of the results. Unlike PCA, t-SNE is a parametric method. The default values for the perplexity, learning rate and number of iterations have been empirically tuned to those that produced acceptable results consistently.

### 5.2.1 Visualizing BERT embeddings with Tensorboard.

Yet another manner to visualize BERT is using TensorBoard, the TensorFlow's open source visualization toolkit which provides all the necessary logic to project the embeddings to a lower dimensional space and to make queries in real time. To that end, we have to write the word embeddings learned weights into disk using a utility described in the `plotter` python module.

```
[5]: # Invoking a method that writes embeddings weights to disk that
# can be loaded with TensorBoard to visualize the embeddings.
out_dir = 'runs/bert_embeddings'
plotter.write_embeddings_to_disk(model,
                                vocab,
                                out_dir=out_dir,
                                write_word_embeddings=True,
                                write_position_embeddings=False,
                                write_type_embeddings=False)
```

You can then run TensorBoard. For example, for `out_dir = 'runs/bert_embeddings'`, you would need to input the following command on a Python console:

```
tensorboard --logdir="<current notebook folder path>\runs"
```

---

## 5.3 Exploring BERT contextual representations.

As previously studied, the power of BERT lies in its ability to derive representations based on context. In this experiment, the embeddings of a word in different contexts (i.e., in different sentences) were computed to check whether there are significant differences among them.

To that end, we used a public domain licensed dataset for word sense disambiguation (WSD) available in [Kaggle](#). The dataset file is an excel file with three columns. The first is the serial number (SN), which assigns a unique identifier to each tuple of contextual sentence (second column) and target polysemous word contained in the sentence (third column).

```
[6]: PATH_WSD_DATASET = 'visualization/test data for WSD evaluation _2905.xlsx'

# Load WSD dataset.
df_polysemy = pd.read_excel(PATH_WSD_DATASET)
```

(continues on next page)

(continued from previous page)

```
# Set the serial number column as the index column.
df_polysemy = df_polysemy.set_index(df_polysemy.sn)

# Let us sample ten random instances.
df_polysemy.sample(n=10, random_state=random_state)
```

```
[6]:
```

	sn	sentence/context	polysemy_word
sn			
583	583	I lost my AC remote.	remote
1812	1812	a young man is running	man
2250	2250	Crane is a large water bird.	crane
1653	1653	Insert the jack in the LAN port	jack
668	668	Thank you for your prompt reply.	prompt
515	515	She likes Russian pop song.	pop
619	619	Films are rated on a scale of poor, fair, good...	scale
2692	2692	The pilot is checking belly of plane	belly
2491	2491	We want to appeal to our core supporters witho...	core
937	937	Any type of single cut metal file can be used ...	file

Let us now use one of the many polysemous words in the dataset.

```
[7]: polysemous_word='bank'
df_polysemy[df_polysemy['polysemy_word']== polysemous_word]
```

```
[7]:
```

	sn	sentence/context	polysemy_word
sn			
1	1	I have bank account.	bank
2	2	Loan amount is approved by the bank.	bank
3	3	He returned to office after he deposited cash ...	bank
4	4	They started using new software in their bank.	bank
5	5	he went to bank balance inquiry.	bank
6	6	I wonder why some bank have more interest rate...	bank
7	7	You have to deposit certain percentage of your...	bank
8	8	He took loan from a Bank.	bank
9	9	he is waking along the river bank.	bank
10	10	The red boat in the bank is already sold.	bank
11	11	Spending time on the bank of Kaligandaki river...	bank
12	12	He was sitting on sea bank with his friend	bank
13	13	She has always dreamed of spending a vacation ...	bank
14	14	Bank of a river is very pleasant place to enjoy.	bank

As it can be seen, the dataset contains fourteen different sentences with the target word bank. Sentences from 1 to 8 refer to bank as as “an organization where people and businesses can invest or borrow money, change it to foreign money, etc., or a building where these services are offered”; whereas sentences from 9 to 14 refer to bank as a “sloping raised land, especially along the sides of a river” (definitions from Cambridge Dictionary). It would be desirable for BERT to be able to disambiguate both senses of the word, which would translate to being capable of deriving distant representations in the n-dimensional space for the different meanings of bank. As we have studied, the self-attention mechanism in BERT’s model architecture is responsible for baking into the representation of each word in a sequence salient information about the rest of the words in the sequence.

This outstanding characteristic of the created representations by BERT exhibit an intrinsic degree of natural language understanding never seen any time before (even at pretraining!). For instance, the context words in the first three sentences of the previous table are different and refer to distinct concepts (e.g., account, loan, amount, deposited, cash). That notwithstanding, they are all related because they all belong to a semantic field of interconnected words that can be used in similar contexts.

Let us test whether BERT realizes this connection and embeds it into the resultant representations.

```
[8]: context_embeddings, labels = plotter.plot_heatmap_embeddings(model,
                                                                tokenizer,
                                                                df_polysemy,
                                                                polysemous_word)
```

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html

Context sentence for each contextual embedding of 'bank'.

```
bank_1: I have bank account.
bank_2: Loan amount is approved by the bank.
bank_3: He returned to office after he deposited cash in the bank.
bank_4: They started using new software in their bank.
bank_5: he went to bank balance inquiry.
bank_6: I wonder why some bank have more interest rate than others.
bank_7: You have to deposit certain percentage of your salary in the bank.
bank_8: He took loan from a Bank.
bank_9: he is waking along the river bank.
bank_10: The red boat in the bank is already sold.
bank_11: Spending time on the bank of Kaligandaki river was his way of enjoying in_
↳his childhood.
bank_12: He was sitting on sea bank with his friend
bank_13: She has always dreamed of spending a vacation on a bank of Caribbean sea.
bank_14: Bank of a river is very pleasant place to enjoy.
```

As it can be seen, BERT seems to, in fact, realize of this connection. Therefore, the embeddings for the target word when referred as a financial institution exhibit strong similarity (the contextual baked into the embedding is similar). Analogously, the embeddings for bank in sentences from 9 to 14 are very similar, and they are all farther from those of sentences from 1 to 8, excepting that of the tenth sentence (arguably because the word sold is in the sentence).

```
[9]: plotter.display_embeddings_scatterplot_3D(context_embeddings,
                                              reduction_strategy='tsne',
                                              random_state=random_state,
                                              _graph_showlegend=False,
                                              title='Projection for different contextual_
↳embeddings',
                                              text=labels,
                                              color=labels)
```

Data type cannot be displayed: application/vnd.plotly.v1+json, text/html



## EXPLORATORY DATA ANALYSIS AND DATA PREPROCESSING.

As part of the undergraduate dissertation: Similarity Measures in Natural Language Processing based on Deep Learning Models.

David Lorenzo Alfaro

### 6.1 Introduction.

To implement a semantic similarity search information retrieval system, a collection of resources is required. We can, then, use similarity measures between searches and data, willing to retrieve relevant information in an efficient fashion.

In our work, we will be using information about ten thousand books from the [goodbooks-10k](#) kaggle dataset. All the information was retrieved from [Goodreads](#), the world's largest site for readers and book recommendations.

The original dataset has been previously modified to better manage the different identifiers and indexes available for each book.

Fortunately, the great majority of word and sentence embedding techniques have been trained on large corpora, often involving humongous number of books (e.g., Toronto Book Corpus). This results on representations that, out of the box, offer great performance for a wide variety of downstream tasks with little fine-tuning being required.

### 6.2 Dataset exploration.

Before getting started, it is first necessary to load all libraries and dependencies that will be used later in the notebook.

```
[4]: import os
import pandas as pd
import numpy as np
from code_utils import utils
from code_utils.paths import *
from pathlib import Path
```

Besides, we set a seed to guarantee reproducibility of the experiments.

```
[5]: seed = 0
```

Let us now unzip containing the collection of books.

```
[6]: from shutil import unpack_archive
unpack_archive(PATH_DATASET_ZIP, DIR_DATASET)
```

```
[7]: filepath = PATH_BOOKS
data = pd.read_csv(filepath, sep=' ')
```

To sample the first  $n$  instances of a dataset we can use the head function.

```
[8]: data.head(5)
```

	book_id	gr_book_id	gr_best_book_id	work_id	books_count	isbn	\
0	0	2767052	2767052	2792775	272	439023483	
1	1	3	3	4640799	491	439554934	
2	2	41865	41865	3212258	226	316015849	
3	3	2657	2657	3275794	487	61120081	
4	4	4671	4671	245494	1356	743273567	

  

	isbn13	authors	original_publication_year	\
0	9.780439e+12	Suzanne Collins	2008.0	
1	9.780440e+12	J.K. Rowling, Mary GrandPré	1997.0	
2	9.780316e+12	Stephenie Meyer	2005.0	
3	9.780061e+12	Harper Lee	1960.0	
4	9.780743e+12	F. Scott Fitzgerald	1925.0	

  

	original_title	... ratings_count	\
0	The Hunger Games	... 4780653	
1	Harry Potter and the Philosopher's Stone	... 4602479	
2	Twilight	... 3866839	
3	To Kill a Mockingbird	... 3198671	
4	The Great Gatsby	... 2683664	

  

	work_ratings_count	work_text_reviews_count	ratings_1	ratings_2	\
0	4942365	155254	66715	127936	
1	4800065	75867	75504	101676	
2	3916824	95009	456191	436802	
3	3340896	72586	60427	117415	
4	2773745	51992	86236	197621	

  

	ratings_3	ratings_4	ratings_5	\
0	560092	1481305	2706317	
1	455024	1156318	3011543	
2	793319	875073	1355439	
3	446835	1001952	1714267	
4	606158	936012	947718	

  

	image_url	\
0	https://images.gr-assets.com/books/1447303603m...	
1	https://images.gr-assets.com/books/1474154022m...	
2	https://images.gr-assets.com/books/1361039443m...	
3	https://images.gr-assets.com/books/1361975680m...	
4	https://images.gr-assets.com/books/1490528560m...	

  

	small_image_url
0	https://images.gr-assets.com/books/1447303603s...
1	https://images.gr-assets.com/books/1474154022s...
2	https://images.gr-assets.com/books/1361039443s...
3	https://images.gr-assets.com/books/1361975680s...
4	https://images.gr-assets.com/books/1490528560s...

[5 rows x 23 columns]

Alternatively, we can use the sample function, which samples  $n$  random instances of the dataset.

```
[9]: data.sample(5, random_state=seed)
```

```
[9]:
```

	book_id	gr_book_id	gr_best_book_id	work_id	books_count	isbn	\
9394	9394	38703	38703	575142	43	385733143	
898	898	53835	53835	1959512	836	159308143X	
2398	2398	43893	43893	1443364	47	765344300	
5906	5906	31244	31244	2888469	378	375761144	
2343	2343	497199	497199	1132770	80	876852630	

  

	isbn13	authors	original_publication_year	\
9394	9.780386e+12	Louis Sachar	2006.0	
898	9.781593e+12	Edith Wharton, Maureen Howard	1920.0	
2398	9.780765e+12	Terry Goodkind	2003.0	
5906	9.780376e+12	Charles Dickens	1865.0	
2343	9.780877e+12	Charles Bukowski	1975.0	

  

	original_title	...	ratings_count	work_ratings_count	\
9394	Small Steps	...	11837	13095	
898	The Age of Innocence	...	102646	114994	
2398	Naked Empire (Sword of Truth, #8)	...	39682	42066	
5906	Our Mutual Friend	...	18599	20659	
2343	Factotum	...	37376	40444	

  

	work_text_reviews_count	ratings_1	ratings_2	ratings_3	ratings_4	\
9394	1387	267	1177	4066	4471	
898	5051	2359	6549	25631	42542	
2398	548	1519	3639	9953	12891	
5906	1102	434	986	3803	6936	
2343	1213	457	1875	8979	16585	

  

	ratings_5	image_url	\
9394	3114	https://s.gr-assets.com/assets/nophoto/book/11...	
898	37913	https://s.gr-assets.com/assets/nophoto/book/11...	
2398	14064	https://s.gr-assets.com/assets/nophoto/book/11...	
5906	8500	https://images.gr-assets.com/books/1403189244m...	
2343	12548	https://images.gr-assets.com/books/1407706616m...	

  

	small_image_url
9394	https://s.gr-assets.com/assets/nophoto/book/50...
898	https://s.gr-assets.com/assets/nophoto/book/50...
2398	https://s.gr-assets.com/assets/nophoto/book/50...
5906	https://images.gr-assets.com/books/1403189244s...
2343	https://images.gr-assets.com/books/1407706616s...

[ 5 rows x 23 columns]

As it can be observed, the dataset has 23 different features. However, only the first and last 10 features of the dataset are being displayed. Let us print the names of all the features in the dataset.

```
[10]: data.columns
```

```
[10]: Index(['book_id', 'gr_book_id', 'gr_best_book_id', 'work_id', 'books_count',  
        'isbn', 'isbn13', 'authors', 'original_publication_year',  
        'original_title', 'title', 'language_code', 'average_rating',  
        'ratings_count', 'work_ratings_count', 'work_text_reviews_count',  
        'ratings_1', 'ratings_2', 'ratings_3', 'ratings_4', 'ratings_5',  
        'image_url', 'small_image_url'],  
        dtype='object')
```

Here's a brief description for the features in the dataset.

- `books_count` represents the number of editions for a given work.
- `gr_best_book_id` contains the most popular edition for a given work.
- Columns `book_id`, `gr_book_id`, `gr_best_book_id`, `work_id`, `isbn` and `isbn13` are different identifiers for the book. As we will see later, the book overviews are not included in this dataset and have been obtained by means of scraping. Each overview is identified with the `gr_book_id` identifier, thus it is the link between both sources of information. Let us first check that it is a valid identifier (i.e., there are no null values and all identifiers are unique).

```
[11]: print(f"Unique values in 'gr_book_id' column: {len(data.gr_book_id.unique())}")
      print(f"Print null values in 'gr_book_id' column {data.gr_book_id[data.gr_book_id.
      ↪isna()]}")
```

```
Unique values in 'gr_book_id' column: 10000
Print null values in 'gr_book_id' column Series([], Name: gr_book_id, dtype: int64)
```

- `gr_book_id` is, in fact, a valid identifier, thus it is the one that we will be using. Furthermore, we also now know that there are no duplicated instances in the dataset, since at least one of its features has no repeated values. The remaining identifier columns can be deleted, as they do not provide any more meaningful information for the tasks that we are to perform.
- As the name suggests, `authors` contains the names of the authors of the book.
- `original_publication_year` indicates the year in which the book was published. We will not be using this information.
- `title` is the english title of the book.
- `original_title` is the title of the book in its original language. We are primarily concerned with english textual information, hence `title` is a more suitable feature.
- `language_code` indicates the textual code assigned to the language of the book. This feature is particularly useful because it will help us get rid of non-English books.
- `average_rating` is a floating value indicating the average rating of a book, ranging from 1 to 5. This feature does not provide relevant information for semantic search, thus it will be discarded. That notwithstanding, it could be used as a criteria to filter the query results, prioritizing those that have better ratings.
- `ratings_count` indicates the number of registered ratings for a book. Analogously, `work_ratings_count` and `work_text_reviews_count` indicate the number of ratings and reviews a work has in the platform, respectively. None of this information is useful for our work.
- `ratings_1`, `ratings_2`, `ratings_3`, `ratings_4` and `ratings_5` characteristics hold the counts for each rating value. Again, this feature does not provide any relevant information to perform semantic search.
- `image_url` and `small_image_url` contain links to pictures of the book cover. Since images cannot be displayed in CLIs, we will discard this information too.

## 6.3 Remove useless features.

Let's get rid of all not useful features.

```
[12]: columns_to_drop = set(['book_id', 'gr_best_book_id', 'work_id', 'books_count',
    'isbn', 'isbn13', 'original_publication_year', 'original_title',
    'average_rating', 'ratings_count', 'work_ratings_count',
    'work_text_reviews_count', 'ratings_1', 'ratings_2', 'ratings_3',
    'ratings_4', 'ratings_5', 'image_url', 'small_image_url'])
```

```
[13]: data = data.drop(columns_to_drop, axis=1)
data.sample(5, random_state=seed)
```

```
[13]:
```

	gr_book_id	authors \		
9394	38703	Louis Sachar		
898	53835	Edith Wharton, Maureen Howard		
2398	43893	Terry Goodkind		
5906	31244	Charles Dickens		
2343	497199	Charles Bukowski		

  

		title	language_code
9394		Small Steps (Holes, #2)	eng
898		The Age of Innocence	eng
2398	Naked Empire (Sword of Truth, #8)		en-GB
5906		Our Mutual Friend	eng
2343		Factotum	NaN

## 6.4 Integrate book overviews into the dataset.

Now that all useless characteristics have been deleted, let's append the overviews to the dataframe. The book overviews are stored in a directory, one txt file for each overview. We will first generate a dataframe containing all txt files in the directory. The filename for each txt file is the gr\_book\_id identifier.

```
[14]: book_overviews = utils.generate_dataframe_from_sparse_txts(DIR_OVERVIEW)
```

```
[15]: print(f"Number of overviews in the dataset: {book_overviews.overview.shape[0]}")
book_overviews
```

```
Number of overviews in the dataset: 9956
```

```
[15]:
```

	gr_book_id	overview
0	1	When Harry Potter and the Half-Blood Prince op...
1	10	Six years of magic, adventure, and mystery mak...
2	10000191	À sa naissance, Lisbeth est enlevée à sa mère ...
3	10006	The discovery of a mysterious notebook turns a...
4	1000751	When orphaned 11-year-old Pollyanna comes to l...
...	...	...
9951	9995135	At long last, New York Times bestselling autho...
9952	99955	Paine's daring prose paved the way for the Dec...
9953	9998	The Woman in the Dunes, by celebrated writer a...
9954	9998705	FLASH! Illuminated by lightning, a lifeless hu...
9955	9999107	Witty, moving, and brilliantly entertaining, T...

```
[9956 rows x 2 columns]
```

As it can be seen, there are 9956 book overviews, which is less than the number of instances in the other dataframe. Consequently, at least 44 books will have no overview. There are several strategies to merge both dataframes. In this case, we will allow having books with no overview (*left join* operation).

```
[16]: data = pd.merge(data, book_overviews, left_on='gr_book_id', right_on='gr_book_id',
↳ how='left')
data
```

```
[16]:
```

	gr_book_id	authors \		title language_code \
0	2767052	Suzanne Collins		
1	3	J.K. Rowling, Mary GrandPré		
2	41865	Stephenie Meyer		
3	2657	Harper Lee		
4	4671	F. Scott Fitzgerald		
...	...	...		
9995	7130616	Ilona Andrews		
9996	208324	Robert A. Caro		
9997	77431	Patrick O'Brian		
9998	8565083	Peggy Orenstein		
9999	8914	John Keegan		
0		The Hunger Games (The Hunger Games, #1)	eng	
1		Harry Potter and the Sorcerer's Stone (Harry P...	eng	
2		Twilight (Twilight, #1)	en-US	
3		To Kill a Mockingbird	eng	
4		The Great Gatsby	eng	
...		...	...	
9995		Bayou Moon (The Edge, #2)	eng	
9996		Means of Ascent (The Years of Lyndon Johnson, #2)	eng	
9997		The Mauritius Command	eng	
9998		Cinderella Ate My Daughter: Dispatches from th...	eng	
9999		The First World War	NaN	
		overview		
0		Winning will make you famous. Losing means cer...		
1		Harry Potter's life is miserable. His parents ...		
2		About three things I was absolutely positive.F...		
3		The unforgettable novel of a childhood in a sl...		
4		On its first publication in 1925, The Great Ga...		
...		...		
9995		The Edge lies between worlds, on the border be...		
9996		Robert A. Caro's life of Lyndon Johnson, which...		
9997		"O'Brian's Aubrey-Maturin volumes actually con...		
9998		The acclaimed author of the groundbreaking bes...		
9999		The First World War created the modern world. ...		

```
[10000 rows x 5 columns]
```

## 6.5 Remove instances with invalid language codes.

Let us now check whether there are noisy data in any of the selected characteristics. Starting off with the language code, we need to make sure that all data fed into the models is in English, being as they have been trained to derive semantic representations for English texts. To that end, let's see how many language codes are in the dataset.

```
[17]: data.language_code.unique()
[17]: array(['eng', 'en-US', 'en-CA', nan, 'spa', 'en-GB', 'fre', 'nl', 'ara',
        'por', 'ger', 'nor', 'jpn', 'en', 'vie', 'ind', 'pol', 'tur',
        'dan', 'fil', 'ita', 'per', 'swe', 'rum', 'mul', 'rus'],
        dtype=object)
```

As it can be seen, there are plenty of different languages. However, is the title and the overview of the book written in the language indicated in language\_code? Let's test it on some of the books with language\_code = spa

```
[18]: data[data.language_code == 'spa'].sample(n=10, random_state=seed)
[18]:
```

	gr_book_id	authors \
9472	53809	Paulo Coelho
83	7677	Michael Crichton
9890	1365225	José Emilio Pacheco
3751	140302	Agatha Christie
4508	63032	Roberto Bolaño
9222	61794	Anonymous
3476	31343	Anne Rice
5125	53926	Mario Vargas Llosa
1799	22590	Philip K. Dick, David Alabort, Manuel Espín
555	10603	Stephen King

  

		title language_code \
9472		Maktub spa
83	Jurassic Park (Jurassic Park, #1)	spa
9890	Las batallas en el desierto	spa
3751	Poirot Investiga (Hércules Poirot, #3)	spa
4508	2666	spa
9222	La vida del Lazarillo de Tormes	spa
3476	Pandora (New Tales of the Vampires, #1)	spa
5125	Travesuras de la niña mala	spa
1799	Ubik	spa
555	Cujo	spa

  

		overview
9472	Maktub não é um livro de conselhos, mas uma tr...	
83		NaN
9890	Historia de un amor imposible, narración de un...	
3751		NaN
4508	A cuatro profesores de literatura, Pelletier, ...	
9222	Lázaro es un muchacho desarrapado a quien la m...	
3476	Anne Rice, creator of the Vampire Lestat, the ...	
5125	¿Cuál es el verdadero rostro del amor?Ricardo ...	
1799	Ubik is a 1969 science fict...	
555	Outside a peaceful town in central Maine, a mo...	

Since the title and the overview seems to be written in the language indicated in language\_code, we will only choose those language codes mapped to English texts: eng, en-US, en-CA, en-GB and en. It is, however, still necessary to check the instances in which the value for the language code is NaN

```
[19]: data[['title', 'overview']][data['language_code'].isna()].sample(n=10, random_
↳state=seed)
```

```
[19]:
          title \
3241 Born Free: A Lioness of Two Worlds (Story of E...
3050 Stone Soup
4807 The Glass Magician (The Paper Magician Trilogy...
9918 Nothing's Fair in Fifth Grade
3971 Experiencing God: Knowing and Doing the Will o...
9772 The Voyages of Doctor Dolittle (Doctor Dolittl...
8179 First Love
2048 Ramona the Pest (Ramona, #2)
9559 Relentless (The Lost Fleet, #5)
9240 Truth Will Prevail (The Work and the Glory, #3)

          overview
3241 There have been many accounts of the return to...
3050 First published in 1947, this classic picture ...
4807 Three months after returning Magician Emery Th...
9918 Jenny knows one thing for sure - Elsie Edwards...
3971 Most Bible studies help people; this one chang...
9772 The delightfully eccentric Doctor Dolittle, re...
8179 An extraordinary portrait of true love that wi...
2048 This is the second title in the hugely popular...
9559 After successfully freeing Alliance POWs, "Bla...
9240 NaN
```

More than 10% of the data has no language code. We verified that all of them are in English, thus they do not have to be deleted. Furthermore, the language\_code feature is no longer needed.

```
[20]: eng_lc = set(['en', 'en-CA', 'en-US', 'en-GB', 'eng'])

data = data[(data.language_code.isin(eng_lc)) | (data.language_code.isna())].drop(
↳'language_code', axis=1)
data.sample(n=10, random_state=seed)
```

```
[20]:
      gr_book_id          authors \
7203      342994 Hans Christian Andersen, Rachel Isadora
8399      53200 Stephen Hawking
8179  17899392 James Patterson, Emily Raymond
7047  2033217 Daniel Silva
1091  17288661 John Grisham
2050      13872 Katherine Dunn
8558  1015311 Ken Akamatsu
6090  21849362 J.R. Ward
6774      7389 Brian K. Vaughan, Adrian Alphona
5760  522525 Carol Tavris, Elliot Aronson

          title \
7203 The Little Match Girl
8399 Black Holes and Baby Universes
8179 First Love
7047 Moscow Rules (Gabriel Allon, #8)
1091 Sycamore Row
2050 Geek Love
8558 Love Hina, Vol. 01
6090 The Shadows (Black Dagger Brotherhood, #13)
6774 Runaways, Vol. 1: Pride and Joy (Runaways, #1)
```

(continues on next page)



(continued from previous page)

```

5760 Mistakes Were Made (But Not by Me): Why We Jus...
                                     overview
7203 The wares of the poor little match girl illumi...
8399 NY Times bestseller. 13 extraordinary essays s...
8179 An extraordinary portrait of true love that wi...
7047 Now the death of a journalist leads Allon to R...
1091 Seth Hubbard is a wealthy man dying of lung ca...
2050 Geek Love is the story of the Binewskis, a car...
8558 At the age of 5, Keitaro and his childhood swe...
6090 Trez "Latimer" doesn't really exist. And not j...
6774 Meet Alex, Karolina, Gert, Chase, Molly and Ni...
5760 Why do people dodge responsibility when things...

```

## 6.6 Remove noisy data from book titles.

The nomenclature utilized for the book titles is as follows: *book\_title + (book\_saga\_name #N\_book\_saga)*. Both the book title and the book saga can be valuable information. However, the book saga number, along with the # symbol may be removed. I have defined a method called `clean_book_title` that allows removing either all saga information or just the saga number.

```
[21]: data.title = [utils.clean_book_title(title) for title in data.title.tolist()]
      data.title.sample(n=10, random_state=seed)
```

```
[21]: 7203                      The Little Match Girl
      8399                  Black Holes and Baby Universes
      8179                      First Love
      7047              Moscow Rules (Gabriel Allon)
      1091                      Sycamore Row
      2050                      Geek Love
      8558              Love Hina, Vol. 01
      6090          The Shadows (Black Dagger Brotherhood)
      6774          Runaways, Vol. 1: Pride and Joy (Runaways)
      5760 Mistakes Were Made (But Not by Me): Why We Jus...
      Name: title, dtype: object
```

## 6.7 Remove noisy data from book overviews.

Luckily, text is automatically tokenized before being fed into any transformer model. That notwithstanding, there is still some work we need to do to clean our text beforehand, like removing special characters, removing extra blank spaces, etc. The `maketrans` built-in method comes handy. It enables us to create a mapping table. We can create an empty mapping table, but the third argument of this function allows us to list all of the characters to remove during the translation process. On the other hand, we will use the `re` module to work with regular expressions with python to further fix some wrong text patterns.

For further details, please check the implementation included in the `utils` module for the `clean_overview` method. Let's see an example:

```
[22]: text = data.overview[data.gr_book_id == 5354].tolist()[0]
      print(f'BEFORE cleaning:\n {text}\n\nAFTER cleaning:\n{utils.clean_overview(text)}')
```

BEFORE cleaning:

Trumble is a minimum-security federal prison, a "camp," home to the usual assortment of relatively harmless criminals--drug dealers, bank robbers, swindlers, embezzlers, tax evaders, two Wall Street crooks, one doctor, at least five lawyers. And three former judges who call themselves the Brethren: one from Texas, one from California, and one from Mississippi. They meet each day in the law library, their turf at Trumble, where they write briefs, handle cases for other inmates, practice law without a license, and sometimes dispense jailhouse justice. And they spend hours writing letters. They are fine-tuning a mail scam, and it's starting to really work. The money is pouring in. Then their little scam goes awry. It ensnares the wrong victim, a powerful man on the outside, a man with dangerous friends, and the Brethren's days of quietly marking time are over.

AFTER cleaning:

Trumble is a minimum-security federal prison, a camp, home to the usual assortment of relatively harmless criminals drug dealers, bank robbers, swindlers, embezzlers, tax evaders, two Wall Street crooks, one doctor, at least five lawyers. And three former judges who call themselves the Brethren: one from Texas, one from California, and one from Mississippi. They meet each day in the law library, their turf at Trumble, where they write briefs, handle cases for other inmates, practice law without a license, and sometimes dispense jailhouse justice. And they spend hours writing letters. They are fine-tuning a mail scam, and it's starting to really work. The money is pouring in. Then their little scam goes awry. It ensnares the wrong victim, a powerful man on the outside, a man with dangerous friends, and the Brethren's days of quietly marking time are over.

```
[23]: data.overview = [utils.clean_overview(str(overview)) for overview in data.overview.  
    ↪ tolist()]
```

Once the preprocessing is done, the dataframe can be exported to a CSV file to avoid repeating these steps everytime we need to work with the cleaned data.

```
[24]: data.set_index('gr_book_id').to_csv(DIR_DATASET + 'books_processed.csv', sep=',')
```

## 6.8 Annex. Code to perform data scraping.

The webpage for each book follow the format [https://www.goodreads.com/book/show/book\\_id](https://www.goodreads.com/book/show/book_id). For instance, <https://www.goodreads.com/book/show/320> is the page containing information for the book “One Hundred Years of Solitude” by Gabriel García Márquez.

The overview is contained in an object called `readable` stacked that can be seen inspecting the code of the page.

```
[ ]: import requests  
    from bs4 import BeautifulSoup  
  
    def scrap_book_overview(book_id, save=False):  
        try:  
            # Connect to the page  
            url = "https://www.goodreads.com/book/show/"+str(book_id)  
            response = requests.get(url)  
            # Instantiate a BeautifulSoup object.  
            soup = BeautifulSoup(response.text, 'lxml')  
            # Access to the component  
            sec = soup.find("div", {"class": "readable stacked"})
```

(continues on next page)

(continued from previous page)

```
# Extract the overview
overview = sec.findAll('span')[-1]
# Store it, should you require it
if not overview.text is None and save:
    file = open("overviews/"+str(book_id)+".txt", "w")
    file.write(overview.text)
    file.close()
    return overview.text
except:
    return None
```



## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## PYTHON MODULE INDEX

### C

`code_utils.plotter`, [17](#)

`code_utils.utils`, [11](#)

### e

`experiment_annoy_ntrees`, [21](#)

`experiment_text_summarization`, [21](#)

### l

`lexical_search`, [7](#)

### s

`semantic_search`, [3](#)





## Symbols

`__init__()` (*lexical\_search.TfidfSearch* method), 7  
`__init__()` (*semantic\_search.SemanticSearch* method), 3

## A

`annoy_index()` (*semantic\_search.SemanticSearch* property), 4  
`append_overviews_to_data()` (in module *code\_utils.utils*), 11  
`apply_dimensionality_reduction()` (in module *code\_utils.plotter*), 17

## B

`biencoder()` (*lexical\_search.TfidfSearch* property), 8

## C

`clean_book_title()` (in module *code\_utils.utils*), 11  
`clean_overview()` (in module *code\_utils.utils*), 12  
`code_utils.plotter` module, 17  
`code_utils.utils` module, 11  
`compute_avg_wordpiece_tokens()` (in module *code\_utils.utils*), 12  
`crossencoder()` (*lexical\_search.TfidfSearch* property), 8  
`crossencoder()` (*semantic\_search.SemanticSearch* property), 4

## D

`display_embeddings_scatterplot_2D()` (in module *code\_utils.plotter*), 17  
`display_embeddings_scatterplot_3D()` (in module *code\_utils.plotter*), 18

## E

`evaluate_n_trees()` (in module *experiment\_annoy\_ntrees*), 21  
`evaluate_summarization_candidates()` (in module *experiment\_text\_summarization*), 21

`experiment_annoy_ntrees` module, 21  
`experiment_text_summarization` module, 21

## F

`fix_punctuation()` (in module *code\_utils.utils*), 12

## G

`generate_dataframe_from_sparse_txts()` (in module *code\_utils.utils*), 12  
`get_bert_model()` (in module *code\_utils.utils*), 13  
`get_bert_tokenizer()` (in module *code\_utils.utils*), 13  
`get_crossencoder()` (in module *code\_utils.utils*), 13  
`get_dir_files_content()` (in module *code\_utils.utils*), 13  
`get_file_id_and_content()` (in module *code\_utils.utils*), 13  
`get_sentence_transformer()` (in module *code\_utils.utils*), 13  
`get_tokenized_text()` (in module *code\_utils.utils*), 14  
`get_top_k_sentences()` (in module *code\_utils.utils*), 14

## H

`histogram_embeddings_nn()` (in module *code\_utils.plotter*), 18

## L

`lexical_search` module, 7  
`load_corpus()` (in module *code\_utils.utils*), 14  
`load_embeddings()` (in module *code\_utils.utils*), 14

## M

`makedirs()` (in module *code\_utils.utils*), 15  
`module`  
`code_utils.plotter`, 17  
`code_utils.utils`, 11

experiment\_annoy\_ntrees, 21  
experiment\_text\_summarization, 21  
lexical\_search, 7  
semantic\_search, 3

## P

plot\_bert\_embeddings\_nn() (in module  
code\_utils.plotter), 18  
plot\_heatmap() (in module code\_utils.plotter), 19  
plot\_heatmap\_embeddings() (in module  
code\_utils.plotter), 19  
plot\_scatter\_with\_secondary\_y\_axis() (in  
module code\_utils.plotter), 19  
prepare\_input\_encoder() (in module  
code\_utils.utils), 15

## R

remove\_filename\_from\_path() (in module  
code\_utils.utils), 15

## S

search() (lexical\_search.TfIdfSearch method), 8  
search() (semantic\_search.SemanticSearch method),  
4  
search\_multiple() (lexical\_search.TfIdfSearch  
method), 8  
search\_multiple() (seman-  
tic\_search.SemanticSearch method), 4  
semantic\_search  
module, 3  
SemanticSearch (class in semantic\_search), 3  
setup\_annoy() (semantic\_search.SemanticSearch  
method), 5  
setup\_bienencoder() (lexical\_search.TfIdfSearch  
method), 8  
setup\_crossencoder() (lexi-  
cal\_search.TfIdfSearch method), 8  
setup\_crossencoder() (seman-  
tic\_search.SemanticSearch method), 5  
split\_text\_into\_sentences\_nltk() (in mod-  
ule code\_utils.utils), 15  
split\_text\_into\_sentences\_spacy() (in  
module code\_utils.utils), 15  
store\_embeddings() (in module code\_utils.utils),  
16  
summarize\_corpus\_overviews() (in module  
code\_utils.utils), 16

## T

test\_annoy\_performance() (seman-  
tic\_search.SemanticSearch method), 5  
TfIdfSearch (class in lexical\_search), 7  
topk\_cos\_sim() (in module code\_utils.utils), 16

## W

write\_embeddings\_to\_disk() (in module  
code\_utils.plotter), 19